

Slide 0 - Introduction

- Only theory questions, e.g. definitions of OS, kernel, etc., the various types of Operating Systems

Slide 1 - Process

- Refer to the [Quiz 1 practice sheet](#) and [solution](#) for practice problems on forking
- Also look at theory questions, e.g. concepts/definitions of process, address space, process states, process control block, system calls, etc.

Slide 2 - Threads

- Theory questions, e.g. definitions, types of threads, threading models, threading issues, etc.
- Questions on the output of threads based on code (must know the behavior of the pthread library and the main functions)
- No need to write an explanation, I am just showing it for clarity.

Examples

1. Find outputs of the following code.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

void *t_func(void *arg);
int var=0;
int t_id[]={1,2};
int main(){
    pthread_t t1;
    pthread_t t2;
    int a1[]={t_id[0],5};
    int a2[]={t_id[1],3};
    pthread_create(&t1,NULL,t_func,(void *)a1);
    pthread_join(t1,NULL);
    pthread_create(&t2,NULL,t_func,(void *)a2);
    pthread_join(t2,NULL);
```

```

    printf("Value of var after operations of threads: %d\n",var);

    return 0;
}
void *t_func(void *arg){
    int *x=arg;
    if(x[0]==1){
        printf("Entered in Thread :%d\n",x[0]);
        var+=x[1];
        printf("Value of var after the operation of Thread %d: %d\n",x[0],var);
        printf("Operation Done by Thread %d...\n",x[0]);
    }
    else{
        printf("Entered in Thread :%d\n",x[0]);
        var-=x[1];
        printf("Value of var after the operation of Thread %d: %d\n",x[0],var);
        printf("Operation Done by Thread %d...\n",x[0]);
    }
}

```

Output	Explanation
Entered in Thread :1	<p>Since pthread_join of t1 was called before thread t2 was created, t1 will always complete entirely before t2 is even started.</p> <p>So, we get the full output of t1, and then we get the full output of t2.</p> <p>Any other ordering for the output will be WRONG for this question.</p>
Value of var after the operation of Thread 1: 5	
Operation Done by Thread 1...	
Entered in Thread :2	
Value of var after the operation of Thread 2: 2	
Operation Done by Thread 2...	
Value of var after operations of threads: 2	

2. Find outputs of the following code.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

void *t_func(void *arg);
int var=0;
int t_id[]={1,2};
int main(){
    pthread_t t1;
    pthread_t t2;
    int a1[]={t_id[0],5};
    int a2[]={t_id[1],3};
    pthread_create(&t1,NULL,t_func,(void *)a1);
    pthread_create(&t2,NULL,t_func,(void *)a2);

    pthread_join(t1,NULL);
    pthread_join(t2,NULL);
    printf("Value of var after operations of threads: %d\n",var);

    return 0;
}
void *t_func(void *arg){
    int *x=arg;
    if(x[0]==1){
        printf("Entered in Thread :%d\n",x[0]);
        var+=x[1];
        printf("Value of var after the operation of Thread %d: %d\n",x[0],var);
        printf("Operation Done by Thread %d...\n",x[0]);
    }
    else{
        printf("Entered in Thread :%d\n",x[0]);
        var-=x[1];
        printf("Value of var after the operation of Thread %d: %d\n",x[0],var);
        printf("Operation Done by Thread %d...\n",x[0]);
    }
}
```

Output

Explanation

Entered in Thread :1	<p>This is a valid output for this one, since both threads were joined after both were created, so there is no definite order for the threads.</p> <p>Sometimes thread 1 finishes first, sometimes thread 2 finishes first, other times both execute at the same time.</p> <p>In this case, any ordering will be valid.</p>
Value of var after the operation of Thread 1: 5	
Operation Done by Thread 1...	
Entered in Thread :2	
Value of var after the operation of Thread 2: 2	
Operation Done by Thread 2...	
Value of var after operations of threads: 2	

Slide 3 - CPU Scheduling

- Given a table of processes, you will have to draw the Gantt chart for the processes. It might be any of the algorithms.
- Theory questions about the different algorithms.
- Refer to the [practice problems for quiz](#) for some problems on CPU Scheduling

Slide 4 - Synchronization

- Theory questions about the fundamentals of synchronization, and the basic structure of locking down a critical section, definitions etc.
- Questions on the output of code with synchronization (must know the behavior of the four methods used for mutex and semaphore, particularly the pthread_mutex library, and the semaphore library)
- No need to write an explanation, I am just showing it for clarity.

Examples

1. Find outputs of the following code.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <semaphore.h>

int t_id[] = {1, 2};
void *t_func1(int *id);
void *t_func2(int *id);
int sum = 15;
pthread_mutex_t m;
sem_t s;

int main() {
    pthread_t t[2];

    sem_init(&s, 0, 0);
    pthread_mutex_init(&m, NULL);

    pthread_create(&t[0], NULL, (void *)t_func1, &t_id[0]);
    pthread_create(&t[1], NULL, (void *)t_func2, &t_id[1]);
```

```

    for (int i = 0; i < 2; i++) {
        pthread_join(t[i], NULL);
    }

    sem_destroy(&s);
    pthread_mutex_destroy(&m);

    printf("Total sum: %d\n", sum);
    return 0;
}

void *t_func1(int *id) {
    sem_wait(&s);

    pthread_mutex_lock(&m);
    for (int i = 0; i < 5; i++) {
        printf("Sum: %d\n", sum);
        sum -= 10;
    }
    pthread_mutex_unlock(&m);
    sem_post(&s);

    return NULL;
}

void *t_func2(int *id) {
    pthread_mutex_lock(&m);
    for (int i = 0; i < 5; i++) {
        printf("Sum: %d\n", sum);
        sum *= 3;
    }
    pthread_mutex_unlock(&m);

    sem_post(&s);
    return NULL;
}

```

Output	Explanation
Sum: 15	In this case, we are using two synchronization tools: - Semaphore: to make sure t_func2 runs before t_func1.
Sum: 45	

Sum: 135	<p>- Mutex: to make sure that the critical session can only be accessed by a single thread.</p> <p>The semaphore is initially set to 0. So if anyone waits for the semaphore then it cannot pass. That is, if t_func1 runs first it will hit the sem_wait and then block there.</p> <p>When t_func2 will run (whether first or second), it doesn't wait, but instead uses sem_post to unlock the semaphore. Then t_func1 can continue.</p> <p>So t_func2 will definitely run before t_func1.</p>
Sum: 405	
Sum: 1215	
Sum: 3645	
Sum: 3635	
Sum: 3625	
Sum: 3615	
Sum: 3605	
Total sum: 3595	

2. Find outputs of the following code.

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>
#include <semaphore.h>

int t_id[] = {1, 2};
void *t_func1(int *id);
void *t_func2(int *id);
int sum = 0;
sem_t s1, s2;

int main() {
    pthread_t t[2];

    sem_init(&s1, 0, 1);
    sem_init(&s2, 0, 0);

    pthread_create(&t[0], NULL, (void *)t_func1, &t_id[0]);
    pthread_create(&t[1], NULL, (void *)t_func2, &t_id[1]);

    for (int i = 0; i < 2; i++) {
        pthread_join(t[i], NULL);
    }
}
```

```

sem_destroy(&s1);
sem_destroy(&s2);

printf("Total sum: %d\n", sum);
return 0;
}

void *t_func1(int *id) {
    sem_wait(&s1);

    for (int i = 0; i < 10; i++) {
        printf("Sum: %d\n", sum);
        sum += 10;
    }

    sem_post(&s1);
    sem_post(&s2);
    return NULL;
}

void *t_func2(int *id) {
    sem_wait(&s2);

    for (int i = 0; i < 10; i++) {
        printf("Sum: %d\n", sum);
        sum -= 5;
    }

    sem_post(&s2);
    return NULL;
}

```

Output	Explanation
Sum: 0 Sum: 10 Sum: 20 Sum: 30 Sum: 40 Sum: 50 Sum: 60 Sum: 70	<p>Same as the last problem, but we are ensuring that t_func1 runs before t_func2, using the semaphore s2.</p> <p>t_func2 waits for s2, which is initially 0. Then t_func1 will use sem_post to unlock s2, and then t_func2 will run.</p>

Sum: 80 Sum: 90 Sum: 100 Sum: 95 Sum: 90 Sum: 85 Sum: 80 Sum: 75 Sum: 70 Sum: 65 Sum: 60 Sum: 55 Total sum: 50	And s1 is used to just behave like a mutex lock.
--	--